

# Landmark Recognition

Deepak Nagaraj, Mike Stackhouse and Roland Lim



# Objective

Recognize well-known landmarks  
based on images with inference  
on the edge: Image Classification


## Why?

- ✓ Organize personal photo collections
- ✓ Visual search
- ✓ Computer vision, autonomous driving
- ✓ Many other commercial applications  
e.g. gaming, education, tourism, \$\$




# Dataset

- The Paris Dataset
- 6412 images collected from Flickr
- 12 Paris Landmarks
  - La Defense Paris
  - Eiffel Tower Paris
  - Hotel des Invalides Paris
  - Louvre Paris
  - Moulin Rouge Paris
  - Musee d'Orsay Paris
  - Notre Dame Paris
  - Pantheon Paris
  - Pompidou Paris
  - Sacre Coeur Paris
  - Arc de Triomphe Paris
  - Paris



# Data preparation - Gathering

- Poor quality images removed
  - Dataset contained quality rating of each image - kept images where  $>25\%$  of object is visible



- Supplement lacking classes
  - Eiffel Tower, Musée d'Orsay, and The Centre Pompidou had  $<100$  images each
  - Used [Google Images Download](#) to add  $\sim 75$  additional images each
- Correct distractor images
  - Images in the “general” class were matched other classes

# Data preparation - Augmentation



- Applied augmentation techniques using the [PIL / Python Image Library](#)



Horizontal Flip



Grayscale



Blur


# Model Development

- Leverage existing models by applying transfer learning
- Took two approaches:
  - NVIDIA DIGITS
    - AlexNet
    - VGG16
  - Tensorflow for Poets
    - MobileNet 0.50
    - Inception V3



# DIGITS

- DIGITS does not have bottlenecking built in but allows you to download a pretrained model
- Trained on TX2 for 3 epochs:
  - AlexNet performance was fairly poor
  - Achieved up to 93% accuracy with VGG16
- VGG16 achieved 97% trained on cloud P100 with 30 epochs
  - Largest concern is that the model is very heavy

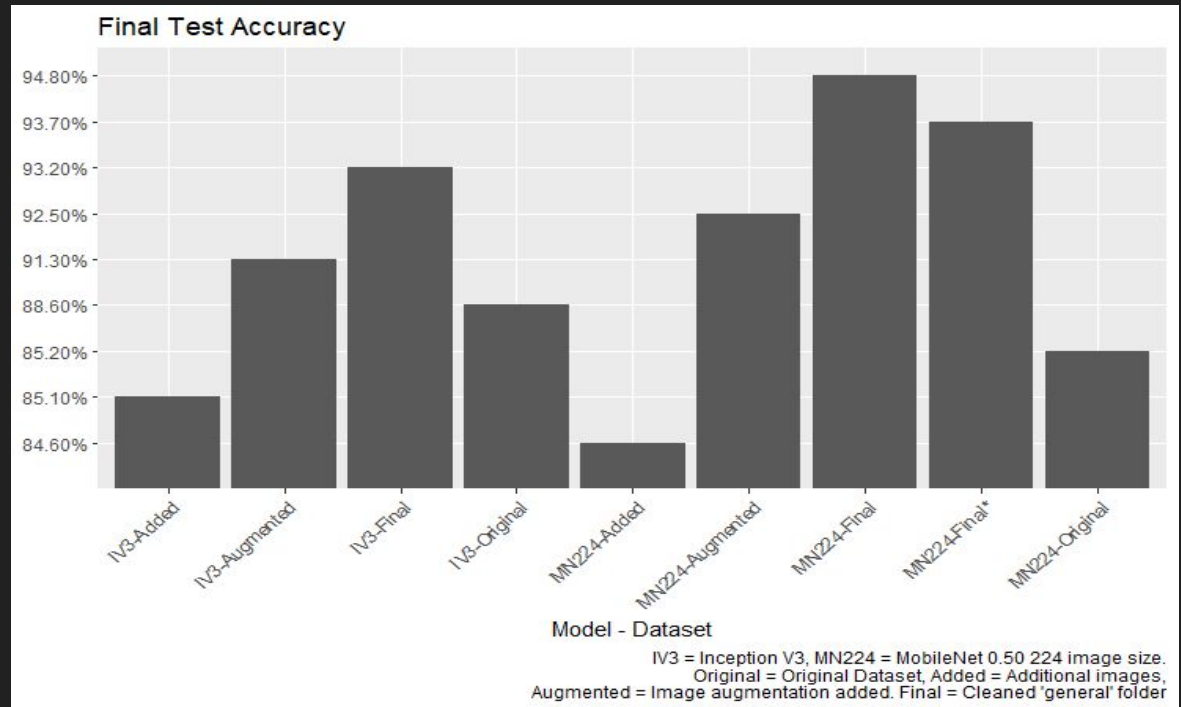


Predictions

defense	98.44%
general	0.71%
pantheon	0.32%
triomphe	0.14%
eiffel	0.13%

# Tensorflow for Poets


- Very efficient training using bottlenecking (<20 minutes)
- Inception V3
  - Top accuracy 93.2%
- MobileNet 0.50
  - Top accuracy 94.8%





# Neural Net Visualization

- Looking inside a neural net to see how image morphs over the layers
- Eiffel Tower over Inception v1:



# Why was that classified wrong?

- Example: Les Invalides misclassified as Sacre Coeur



# Conclusion and Next Steps

- MobileNet is our ultimate architecture of choice, despite higher performance from VGG16 trained on a cloud P100
  - MobileNet architecture is much smaller and more practical for edge computation, yet still achieved 94.8% accuracy
- Next steps would be to expand the number of landmarks classified
  - Need a larger dataset (i.e. Google Landmark Recognition Challenge)
  - Image generalization could become problematic, so possibly link location data
- Could apply model pruning leveraging NVIDIA transfer learning toolkit
  - Cuts model size down further, faster computation

# References

- <http://www.robots.ox.ac.uk/~vgg/data/parisbuildings/>
- <https://www.kaggle.com/c/landmark-recognition-challenge>
- <https://landmarksworkshop.github.io/CVPRW2018/>
- <https://www.blippar.com/blog/2018/02/16/landmark-recognition-api-never-confuse-landmark-again>
- <https://cloud.google.com/blog/products/gcp/around-the-world-landmark-detection-with-the-cloud-vision-api>
- <https://medium.com/@abhinaya08/google-landmark-recognition-274aab3c71ae>
- <https://towardsdatascience.com/google-landmark-recognition-using-transfer-learning-dde35cc760e1>
- <https://medium.com/@abhinaya08/google-landmark-recognition-274aab3c71ae>
- Lucid Visualization Library:  
[https://colab.research.google.com/github/tensorflow/lucid/blob/master/notebooks/misc/feature\\_inversion\\_caricatures.ipynb](https://colab.research.google.com/github/tensorflow/lucid/blob/master/notebooks/misc/feature_inversion_caricatures.ipynb)